

## Ambiente Docker levantado com sucesso pelo:

- [https://github.com/fabiogjardim/bigdata\\_docker.git](https://github.com/fabiogjardim/bigdata_docker.git)

Alguns exemplos das aplicações rodando,

The screenshot shows the Docker Desktop interface. On the left, there's a sidebar with options like Ask Gordon, Containers (selected), Images, Volumes, Builds, Models (Beta), and MCP Toolkit (Beta). The main area is titled 'Containers' with a sub-section 'Container CPU usage'. It shows 97.76% / 600% (6 CPUs available) and 3.94GB / 15.17GB memory usage. A search bar and a filter 'Only show running containers' are at the top. Below is a table of running containers:

Name	Container ID	Image	Port(s)	CPU (%)	Last started	Actions
bigdata_docker	-	-	-	97.34%	3 seconds ago	[...]
presto	c708c4b92a07	fjardim/presto	8080:8080	2.78%	6 minutes ago	[...]
hive-server	4c6dca88b5de	fjardim/hive	10000:10000	0.54%	6 minutes ago	[...]
hive_metastore	356917fb38e2	fjardim/hive	9083:9083	0.48%	6 minutes ago	[...]
kafkamanager	52fa85ec5665	fjardim/kafkamanager	9000:9000	0%	7 seconds ago	[...]
hive-metastore-postgresql	b3a852867358	fjardim/hive-metastore	-	0%	6 minutes ago	[...]
hue	e0374789b1da	fjardim/hue	8888:8888	0.04%	6 minutes ago	[...]
hbase-master	39916576f516	fjardim/hbase-master	16010:16010	0.24%	6 minutes ago	[...]
kafka	52a3dc856e45	fjardim/kafka	9092:9092	0.44%	6 minutes ago	[...]
datanode	3cd091585c16	fjardim/datanode	50075:50075	0.19%	6 minutes ago	[...]
mongo	d51a61e3c361	fjardim/mongo	22017:22017	0.2%	6 minutes ago	[...]
database	a9edfed1ecb	fjardim/mysql	33061:3306	0.02%	6 minutes ago	[...]
nifi	c1824c5bb982	fjardim/nifi	8443:8443	91.25%	4 seconds ago	[...]

At the bottom, it says 'Showing 19 items'.

```
Windows PowerShell
Copyright (C) Microsoft Corporation. Todos os direitos reservados.

Instale o PowerShell mais recente para obter novos recursos e aprimoramentos! https://aka.ms/PSWindows

PS C:\Users\Willgnner\Documents\bigdata_docker> docker-compose up -d
time="2025-12-14T21:34:20-03:00" level=warning msg="C:\\\\Users\\\\Willgnner\\\\Documents\\\\bigdata_docker\\\\docker-compose.yml: the attribute 'version' is obsolete, it will be ignored, please remove it to avoid potential confusion"
[+] Running 17/17
  ✓ Container mongo_express          Started      0.4s
  ✓ Container jupyter-spark          Started      0.7s
  ✓ Container metastable            Started      0.5s
  ✓ Container database              Started      0.8s
  ✓ Container zookeeper             Started      0.4s
  ✓ Container namenode              Started      0.6s
  ✓ Container mongo                 Started      0.9s
  ✓ Container nifi                  Started      0.8s
  ✓ Container kafka                 Started      0.6s
  ✓ Container hbase-master          Started      0.7s
  ✓ Container datanode              Started      0.6s
  ✓ Container hue                   Started      0.9s
  ✓ Container kafkamanager          Started      0.7s
  ✓ Container hive-metastore-postgresql Started      0.6s
  ✓ Container hive-metastore        Started      0.6s
  ✓ Container hive-server           Started      0.6s
  ✓ Container presto                Started      0.7s
PS C:\Users\Willgnner\Documents\bigdata_docker> |
```

**HDFS:**

[localhost:50070/dshealth.html#tab-overview](http://localhost:50070/dshealth.html#tab-overview)

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities -

### Overview 'namenode:8020' (active)

Started:	Mon Dec 15 00:34:29 UTC 2025
Version:	2.7.4, rcc915e6b9d0131462a07301566c175728a262
Compiled:	2017-08-01T00:29Z by kslivachik from branch-2.7.4
Cluster ID:	CID-c2019425-f9bc-4bb9-ad29-c36a0d369a6c
Block Pool ID:	BP-1660274043-172.18.0.2-1765141395281

### Summary

Security is off.  
Safemode is off.  
269 files and directories, 213 blocks = 482 total filesystem object(s).  
Heap Memory used 142.55 MB of 244 MB Heap Memory. Max Heap Memory is 889 MB.  
Non Heap Memory used 40.82 MB of 41.88 MB Committed Non Heap Memory. Max Non Heap Memory is -1 B.

Configured Capacity:	465.02 GB
DFS Used:	114.26 MB (0.02%)
Non DFS Used:	447.3 GB
DFS Remaining:	17.61 GB (3.79%)
Block Pool Used:	114.26 MB (0.02%)
DataNodes usages% (Min/Median/Max/stdDev):	0.02% / 0.02% / 0.02% / 0.00%
Live Nodes	1 (Decommissioned: 0)
Dead Nodes	0 (Decommissioned: 0)

## Presto:

[localhost:8080](http://localhost:8080)

### CLUSTER OVERVIEW

VERSION: 0.181 ENVIRONMENT: DEVELOPMENT UPTIME: 1.78m

RUNNING QUERIES: 0	ACTIVE WORKERS: 1	ROWS/SEC: 0.00
QUEUED QUERIES: 0	RUNNABLE DRIVERS: 0.00	BYTES/SEC: 0
BLOCKED QUERIES: 0	RESERVED MEMORY (B): 0	WORKER PARALLELISM: 0.00

### QUERY DETAILS

User, source, Query ID or query text	Filter	Running/blocked	Queued	Finished	Failed	User error	Sort	Reorder Interval	Show
No queries									

## Hbase:

[localhost:16010/master-status](http://localhost:16010/master-status)

**Master** hbase-master

### Region Servers

### Backup Masters

ServerName	Port	Start Time
Total 0		

### Tables

- User Tables
- System Tables
- Snapshots

### Tasks

- Show All Monitored Tasks
- Show non-RPC Tasks
- Show All RPC Handler Tasks
- Show Active RPC Calls
- Show Client Operations
- Show as JSON

Start Time	Description	State	Status
Mon Dec 15 00:34:39 UTC 2025	Master startup	RUNNING (since 43sec ago)	Waiting for region servers count to settle, currently checked in 0, slept for 1523 ms, expecting minimum of 1, maximum of 2147483647, timeout of 4500 ms, interval of 1500 ms. (since 0sec ago)

### Software Attributes

Attribute Name	Value	Description
HBase Version	1.2.6, revision=Unknown	HBase version and revision

## Mongo Express:

The screenshot shows the Mongo Express web application running at localhost:8081. At the top, there's a navigation bar with a magnifying glass icon, a search bar containing "localhost:8081", and several browser control icons. Below the header, the title "Mongo Express" is displayed. The main content area is divided into two sections: "Databases" and "Server Status".

**Databases:** This section lists four databases: admin, config, local, and pmdb. Each database entry has a green "View" button on the left and a red "Del" button on the right. A "Create Database" button is located at the top right of this section.

Databases			
	admin		Del
	config		Del
	local		Del
	pmdb		Del

**Server Status:** This section provides various metrics about the MongoDB server. It includes tables for Hostname, Uptime, Current Connections, Active Clients, Clients Reading, Read Lock Queue, Disk Flushes, and Time Spent Flushing.

Server Status			
Hostname	mongo	MongoDB Version	4.2.3
Uptime	62 seconds	Server Time	Mon, 15 Dec 2025 00:35:23 GMT
Current Connections	4	Available Connections	838856
Active Clients	0	Queued Operations	0
Clients Reading	0	Clients Writing	0
Read Lock Queue	0	Write Lock Queue	0
Disk Flushes		Last Flush	
Time Spent Flushing	ms	Average Flush Time	ms

## Metabase:

The screenshot shows the Metabase interface. At the top, there's a blue header bar with a search bar containing "Q. Pesarlar...", a user profile icon, and several other navigation icons. The main content area is divided into several sections.

**Dashboard Creation:** On the left, there's a sidebar with a message "Como vai, bigdata123?" and a "COMECE AQUI" button. Below it, a message says "Os painéis mais importantes do seu time vão aqui" and "Fixe painéis em Nossas análises para eles aparecerem neste espaço para todos".

**Analyses:** A section titled "EXPERIMENTE ESSES RAIOS-X QUE CRIEI COM SEUS DADOS." shows a placeholder for "Nossas análises" with a "Salvar painéis, questões e coleções em 'Nossas análises'" button and a "Exibir todos os itens >" link.

**Data Sources:** A section titled "NOSSOS DADOS" lists three data sources: "mongodb" (with a purple cylinder icon), "Presto" (with a purple cylinder icon), and "Sample Dataset" (with a purple cylinder icon).

## Jupyter Spark:



## Hue:

The screenshot shows the Hue interface with a query history and a table list. The query history contains numerous identical SQL statements for creating external tables and dropping existing ones. The table list on the right shows 'No tables identified'.

## Máquina Virtual

The screenshot shows the Oracle VM VirtualBox Manager interface. A window titled 'ubuntu-server-18.04-hdp-itr [Executando]' is open, showing the terminal output of a Jupyter notebook. The configuration pane on the left shows the virtual machine's settings, including its name ('ubuntu-server-18.04-hdp-itr'), processor ('Intel Dual Band'), memory ('8192 MB'), and storage ('Hortonworks Docker Sandbox HDP'). The terminal window shows the user logging in and starting a Jupyter notebook.

## Jupyter:

A screenshot of a web browser displaying a file tree titled "jupyter". The tree shows the contents of the root directory, including "common", "elf", "files", "ml", "delta-start.ipynb", "config.yaml", and "start-jupyter.sh". Each item has a checkbox next to it. At the top right are "Upload" and "New" buttons, and at the bottom right are "Quit" and "Logout" buttons.

## Ambari:

A screenshot of the Ambari "Background Operations" dialog. It lists six operations under "1 Background Operation Running": "Start All Services" (status 100%, user admin, start time Today 13:36, duration 3m 57s), "Stop All Services" (status 100%, user admin, start time Mon Aug 18 2025 14:10, duration 47s), "Start All Services" (status 100%, user admin, start time Mon Aug 18 2025 14:05, duration 4m 17s), "Stop All Services" (status 100%, user admin, start time Tue Nov 05 2024 14:04, duration 46s), "Start All Services" (status 100%, user admin, start time Tue Nov 05 2024 14:01, duration 2m 55s), and "Stop All Services" (status 100%, user admin, start time Sun Oct 13 2024 19:17, duration 1m 2s). A checkbox at the bottom left says "Do not show this dialog again when starting a background operation".

A screenshot of the Ambari Metrics dashboard. The top navigation bar shows "localhost:8080/#/main/dashboard/metrics" and "Dashboard / Metrics". The top right shows user information: "bigdata", a gear icon, a bell icon with 3 notifications, a grid icon, and "admin". Below the navigation is a toolbar with "METRICS" (selected), "HEATMAPS", "CONFIG HISTORY", "METRIC ACTIONS", and "LAST 1 HOUR". The main area contains a 4x2 grid of metrics: NameNode Heap (6%), HDFS Disk Usage (81%), NameNode CPU WIO (n/a), DataNodes Live (1/1); NameNode RPC (0 ms), Memory Usage (No Data Available), Network Usage (No Data Available), CPU Usage (No Data Available); Cluster Load (No Data Available), NameNode Uptime (4m 25s), ResourceManager Heap (15%), NodeManagers Live (1/1). A message at the bottom right says "Ativar o Windows" and "Acesse Configurações para ativar o Windows."



Ambari

Dashboard

Services

... ▾

● HDFS 2

● YARN 1

● MapReduce2

● ZooKeeper

● Ambari Metrics 1

● Kafka

● SmartSense 1

Hosts